

## Artificial Intelligence Questions

Question 1: What is the primary goal of Artificial Intelligence (AI)?

- a) To simulate human emotions
- b) To build intelligent machines that mimic human intelligence
- c) To create robots with physical abilities
- d) To replace human creativity

Answer: b) To build intelligent machines that mimic human intelligence

Question 2: Which of the following is an example of Weak AI (Narrow AI)?

- a) A chatbot that can have human-like conversations
- b) A robot that can perform complex tasks autonomously
- c) A self-driving car that can navigate in real-world traffic
- d) A chess-playing program that can defeat human players

Answer: d) A chess-playing program that can defeat human players

Question 3: What is the main challenge in achieving Strong AI (Artificial General Intelligence)?

- a) Building faster processors
- b) Creating emotional machines
- c) Replicating human consciousness
- d) Understanding human creativity

Answer: c) Replicating human consciousness

Question 4: Which of the following is NOT a subfield of Artificial Intelligence?

- a) Machine Learning
- b) Robotics
- c) Data Wrangling
- d) Natural Language Processing

Answer: c) Data Wrangling

Question 5: What is Python?

- a) A high-level programming language
  - b) A low-level programming language
  - c) A database management system
  - d) A graphics design software
- Answer: a) A high-level programming language

Question 6: Which keyword is used to define a function in Python?

- a) func
- b) function
- c) def
- d) define

Answer: c) def

Question 7: What is the output of the following Python code snippet?

```
python
Copy code
x = [1, 2, 3]
y = x
x[0] = 4
```

print(y)

- a) [1, 2, 3]
- b) [4, 2, 3]
- c) [1, 2, 3, 4]
- d) [4, 2, 3, 4]

Answer: d) [4, 2, 3]

Question 8: Which of the following is used to represent a single-line comment in Python?

- a) /\* ... \*/
- b) <!-- ... -->
- c) # ...
- d) // ...

Answer: c) # ...

Question 9: What is data wrangling?

- a) A process of data visualization
- b) A process of data preprocessing and cleaning
- c) A process of data analysis
- d) A process of data model training

Answer: b) A process of data preprocessing and cleaning

Question 10: Which of the following is NOT a common data wrangling task?

- a) Data cleaning
- b) Data visualization
- c) Data transformation
- d) Data imputation

Answer: b) Data visualization

Question 11: Which Python library is commonly used for data wrangling tasks?

- a) Pandas
- b) Numpy
- c) Scikit-learn
- d) Matplotlib

Answer: a) Pandas

Question 12: What does the "NaN" value represent in Pandas?

- a) It stands for "Not a Name" and indicates missing data.
- b) It stands for "Numerical and Numeric" and represents numeric values.
- c) It stands for "None and Null" and represents missing data.
- d) It stands for "No Action Necessary" and indicates complete data.

Answer: c) It stands for "None and Null" and represents missing data.

Question 13: What is the basic building block of a neural network?

- a) Neuron
- b) Layer
- c) Node
- d) Weight

Answer: a) Neuron

Question 14: In a neural network, what is the activation function responsible for?

- a) Calculating the weighted sum of inputs
- b) Predicting the output value
- c) Introducing non-linearity to the model

d) Adjusting the weights during training

Answer: c) Introducing non-linearity to the model

Question 15: What is the term used for the process of fine-tuning the neural network's parameters to minimize errors during training?

a) Forward Propagation

b) Backward Propagation

c) Gradient Descent

d) Model Optimization

Answer: b) Backward Propagation

Question 16: Which type of neural network is used for problems with sequential or time-series data?

a) Feedforward Neural Network (FNN)

b) Convolutional Neural Network (CNN)

c) Recurrent Neural Network (RNN)

d) Multilayer Perceptron (MLP)

Answer: c) Recurrent Neural Network (RNN)

Question 17: Which popular deep learning framework is developed by Google and widely used for building neural networks?

a) PyTorch

b) Tensorflow

c) Keras

d) Theano

Answer: b) Tensorflow

Question 18: What is the purpose of Keras in deep learning?

a) To handle image data in deep learning tasks

b) To preprocess and clean the data before training

c) To build and train neural networks with a user-friendly interface

d) To perform hyperparameter optimization during training

Answer: c) To build and train neural networks with a user-friendly interface

Question 19: In Keras, what does the number of nodes in the input layer depend on?

a) The number of hidden layers

b) The size of the training dataset

c) The number of features in the input data

d) The learning rate during training

Answer: c) The number of features in the input data

Question 20: Which function in Keras allows you to compile a neural network with a specified loss function and optimizer?

a) model.summary()

b) model.compile()

c) model.fit()

d) model.predict()

Answer: b) model.compile()

Question 21: What is the main advantage of using Convolutional Neural Networks (CNNs) for image recognition tasks?

a) Faster training time compared to traditional neural networks

b) Ability to handle text data effectively

c) Ability to capture spatial relationships in the image

d) Higher accuracy on numerical datasets

Answer: c) Ability to capture spatial relationships in the image

Question 22: Which layer in a CNN is responsible for reducing the spatial dimensions of the input image?

a) Convolutional Layer

b) Pooling Layer

c) Dense Layer

d) Activation Layer

Answer: b) Pooling Layer

Question 23: In a CNN, what are the learnable filters used for?

a) They detect edges and textures in the input image.

b) They determine the size of the output feature map.

c) They represent the number of hidden neurons in the network.

d) They initialize the weights during training.

Answer: a) They detect edges and textures in the input image.

Question 24: Which activation function is commonly used in the convolutional layers of a CNN?

a) Sigmoid

b) ReLU (Rectified Linear Activation)

c) Tanh (Hyperbolic Tangent)

d) Softmax

Answer: b) ReLU (Rectified Linear Activation)

Question 25: What is the key characteristic of a Long Short-Term Memory (LSTM) cell in an RNN?

a) It can process sequences in parallel.

b) It can only handle short sequences of data.

c) It uses a single layer for sequential processing.

d) It can retain information over long time intervals.

Answer: d) It can retain information over long time intervals.

Question 26: What is the primary purpose of the "embedding" layer in an RNN for Natural Language Processing tasks?

a) To convert text data into numerical representations.

b) To handle sequential data with time dependencies.

c) To reduce the dimensionality of the input data.

d) To apply non-linear activation functions to the input data.

Answer: a) To convert text data into numerical representations.

Question 27: What is the primary challenge of using traditional RNNs in long sequences?

a) They require more computational resources.

b) They are prone to overfitting the data.

c) They suffer from the vanishing gradient problem.

d) They have difficulty handling short sequences.

Answer: c) They suffer from the vanishing gradient problem.

Question 28: Which RNN variant is designed to solve the vanishing gradient problem by using gated cells?

a) Feedforward Neural Network (FNN)

b) Long Short-Term Memory (LSTM)

c) Elman Recurrent Neural Network (RNN)

d) Radial Basis Function (RBF) Neural Network

Answer: b) Long Short-Term Memory (LSTM)

Question 29: What is Natural Language Processing (NLP)?

- a) A technique to convert numerical data into natural language text.
- b) A process of extracting features from images using deep learning.
- c) A field that focuses on the interaction between computers and human language.
- d) A method to encode text data into binary format for machine learning models.

Answer: c) A field that focuses on the interaction between computers and human language.

Question 30: Which NLP task involves classifying text documents into predefined categories?

- a) Sentiment Analysis
- b) Named Entity Recognition (NER)
- c) Topic Modeling
- d) Text Classification

Answer: d) Text Classification

Question 31: What is the purpose of tokenization in NLP?

- a) To convert text into lowercase letters only.
- b) To split text into individual words or tokens.
- c) To remove stop words from the text.
- d) To convert text into its numerical representation.

Answer: b) To split text into individual words or tokens.

Question 32: Which Python library is commonly used for NLP tasks, such as tokenization, stemming, and lemmatization?

- a) TensorFlow
- b) PyTorch
- c) NLTK (Natural Language Toolkit)
- d) Scikit-learn

Answer: c) NLTK (Natural Language Toolkit)

Question 33: Which IBM Cloud service provides pre-trained machine learning models that can be easily deployed for various tasks?

- a) IBM Watson Assistant
- b) IBM Watson Studio
- c) IBM Watson Discovery
- d) IBM Watson Machine Learning

Answer: d) IBM Watson Machine Learning

Question 34: Which IBM Cloud service allows users to build, train, and deploy AI models without requiring extensive coding knowledge?

- a) IBM Watson Assistant
- b) IBM Watson Studio
- c) IBM Watson Discovery
- d) IBM Watson Natural Language Understanding

Answer: b) IBM Watson Studio

Question 35: What is the primary purpose of deploying a machine learning model on the IBM Cloud?

- a) To create visualizations of the model's predictions
- b) To convert the model into a Python script
- c) To make predictions on new data in a production environment
- d) To handle data preprocessing before model deployment

Answer: c) To make predictions on new data in a production environment

Question 36: Which IBM Cloud service provides a cloud-based platform for building, training, and deploying machine learning models?

- a) IBM Watson Assistant
  - b) IBM Watson Studio
  - c) IBM Watson Discovery
  - d) IBM Watson Text to Speech
- Answer: b) IBM Watson Studio

Question 37: What is the primary goal of building and deploying a machine learning application?

- a) To achieve the highest possible accuracy during model training
  - b) To visualize data and insights obtained from the model
  - c) To make the model accessible and usable for end-users
  - d) To create complex machine learning models using deep learning techniques
- Answer: c) To make the model accessible and usable for end-users

Question 38: Which cloud platform allows you to deploy machine learning models as web applications?

- a) AWS (Amazon Web Services)
  - b) Google Cloud Platform
  - c) Microsoft Azure
  - d) IBM Cloud
- Answer: d) IBM Cloud

Question 39: What is the role of API (Application Programming Interface) in deploying a machine learning model?

- a) To convert the model into a human-readable format
  - b) To enable communication between the model and other software applications
  - c) To generate visualizations of the model's predictions
  - d) To handle data preprocessing before model deployment
- Answer: b) To enable communication between the model and other software applications

Question 40: Which deployment method allows you to execute machine learning models on mobile devices?

- a) On-Premise Deployment
  - b) Cloud Deployment
  - c) Edge Deployment
  - d) Web Deployment
- Answer: c) Edge Deployment

Question 41: What is the primary goal of data science?

- a) To study the history of data
  - b) To analyze and interpret data to gain insights
  - c) To create complex mathematical models
  - d) To develop software applications
- Answer: b) To analyze and interpret data to gain insights

Question 42: Which step in the data science process involves defining the problem and setting the research objectives?

- a) Data Collection
  - b) Data Analysis
  - c) Data Cleaning
  - d) Problem Formulation
- Answer: d) Problem Formulation

Question 43: Which of the following is NOT a valid Python data type?

- a) Integer
- b) String

- c) Dictionary
- d) Loop

Answer: d) Loop

Question 44: What is the output of the following Python code snippet?

```
python
Copy code
x = 5
y = 2
z = x % y
print(z)
```

- a) 2.5
- b) 3
- c) 0.5
- d) 1

Answer: d) 1

Question 45: Data wrangling is also known as:

- a) Data Exploration
- b) Data Preprocessing
- c) Data Visualization
- d) Data Modeling

Answer: b) Data Preprocessing

Question 46: Which data wrangling technique involves filling missing values with the mean or median of the existing data?

- a) Data Imputation
- b) Data Scaling
- c) Data Normalization
- d) Data Encoding

Answer: a) Data Imputation

Question 47: In regression, the dependent variable is:

- a) Categorical
- b) Continuous
- c) Discrete
- d) Binary

Answer: b) Continuous

Question 48: What is the purpose of the coefficient of determination (R-squared) in regression?

- a) To measure the variance of the independent variable
- b) To measure the correlation between independent and dependent variables
- c) To assess the goodness of fit of the regression model
- d) To calculate the standard error of the regression coefficients

Answer: c) To assess the goodness of fit of the regression model

Question 49: Which algorithm is commonly used for binary classification tasks?

- a) k-Nearest Neighbors (k-NN)
- b) Random Forest
- c) Support Vector Machine (SVM)
- d) K-Means Clustering

Answer: c) Support Vector Machine (SVM)

Question 50: What is the output of the logistic function (sigmoid) used in logistic regression?

- a) Positive integers
- b) Negative integers
- c) Real numbers between 0 and 1
- d) Binary values (0 or 1)

Answer: c) Real numbers between 0 and 1

Question 51: What is the role of hyperparameters in machine learning algorithms?

- a) To control the model complexity
- b) To represent the target variable
- c) To identify outliers in the data
- d) To perform feature selection

Answer: a) To control the model complexity

Question 52: Which technique aims to find the optimal hyperparameters by trying different combinations exhaustively?

- a) Grid Search
- b) Random Search
- c) Gradient Descent
- d) K-Means Clustering

Answer: a) Grid Search

Question 53: What is the primary goal of unsupervised learning?

- a) To make predictions based on historical data
- b) To classify data into different categories
- c) To find patterns and structures in the data
- d) To measure the performance of a model

Answer: c) To find patterns and structures in the data

Question 54: Which unsupervised learning algorithm is commonly used for dimensionality reduction?

- a) Decision Tree
- b) Principal Component Analysis (PCA)
- c) Naive Bayes
- d) Linear Regression

Answer: b) Principal Component Analysis (PCA)

Question 55: In a binary classification problem, which metric is used to evaluate the model's performance when correctly identifying positive cases?

- a) Precision
- b) Recall
- c) F1-score
- d) Accuracy

Answer: b) Recall

Question 56: What is the formula for calculating the F1-score using precision and recall?

- a)  $F1\text{-score} = (\text{Precision} + \text{Recall}) / 2$
- b)  $F1\text{-score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$
- c)  $F1\text{-score} = \text{Precision} - \text{Recall}$
- d)  $F1\text{-score} = \text{Precision} * \text{Recall}$

Answer: b)  $F1\text{-score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$

Question 57: What is the purpose of IBM Watson Assistant?



- a) To analyze and visualize data
  - b) To build and deploy machine learning models
  - c) To provide a conversational interface for applications
  - d) To perform natural language processing tasks
- Answer: c) To provide a conversational interface for applications

Question 58: Which IBM Watson service is used to analyze unstructured text data and extract insights from it?

- a) IBM Watson Assistant
  - b) IBM Watson Discovery
  - c) IBM Watson Studio
  - d) IBM Watson Machine Learning
- Answer: b) IBM Watson Discovery

Question 59: What is the role of a REST API in machine learning model deployment?

- a) To convert the model into a human-readable format
  - b) To handle data preprocessing before model deployment
  - c) To enable communication between the model and other software applications
  - d) To generate visualizations of the model's predictions
- Answer: c) To enable communication between the model and other software applications

Question 60: Which deployment method is suitable for machine learning models that require real-time predictions and low latency?

- a) Cloud Deployment
  - b) On-Premise Deployment
  - c) Edge Deployment
  - d) Web Deployment
- Answer: c) Edge Deployment

Question 61: Which of the following is an example of Strong AI (Artificial General Intelligence)?

- a) A chatbot that can have human-like conversations
  - b) A robot that can perform complex tasks autonomously
  - c) A self-driving car that can navigate in real-world traffic
  - d) A computer program that can pass the Turing Test
- Answer: d) A computer program that can pass the Turing Test

Question 62: What is the main difference between Artificial Intelligence and Machine Learning?

- a) Artificial Intelligence involves making machines intelligent, while Machine Learning involves training models to perform tasks.
  - b) Artificial Intelligence is a subset of Machine Learning.
  - c) Artificial Intelligence is used for creating robots, while Machine Learning is used for natural language processing.
  - d) Artificial Intelligence is a general term, while Machine Learning is a specific technique.
- Answer: a) Artificial Intelligence involves making machines intelligent, while Machine Learning involves training models to perform tasks.

Question 63: Which data structure in Python is used to store a collection of elements that are unordered and unchangeable?

- a) List
  - b) Tuple
  - c) Dictionary
  - d) Set
- Answer: d) Set

Question 64: What is the output of the following Python code snippet?

python

Copy code

```
x = "Hello"
```

```
y = "World"
```

```
z = x + y
```

```
print(z)
```

a) Hello

b) World

c) HelloWorld

d) HeloWrld

Answer: c) HelloWorld

Question 65: What is the purpose of data normalization in data wrangling?

a) To remove outliers from the dataset

b) To handle missing data

c) To scale numerical features to a similar range

d) To encode categorical variables

Answer: c) To scale numerical features to a similar range

Question 66: Which data wrangling technique is used to transform categorical variables into numerical representations?

a) Data Imputation

b) Data Normalization

c) Data Encoding

d) Data Scaling

Answer: c) Data Encoding

Question 67: What is the purpose of the activation function in a neural network?

a) To calculate the weighted sum of inputs

b) To determine the number of hidden layers

c) To introduce non-linearity to the model

d) To adjust the weights during training

Answer: c) To introduce non-linearity to the model

Question 68: In a neural network, what does the term "weights" refer to?

a) The input features of the model

b) The output labels of the model

c) The bias terms used in the model

d) The adjustable parameters that determine the strength of connections between neurons

Answer: d) The adjustable parameters that determine the strength of connections between neurons

Question 69: Which deep learning framework is developed by Facebook's AI Research (FAIR) lab?

a) PyTorch

b) Tensorflow

c) Keras

d) Caffe

Answer: a) PyTorch

Question 70: What is the purpose of the activation function in a neural network?

a) To calculate the weighted sum of inputs

b) To determine the number of hidden layers

- c) To introduce non-linearity to the model
- d) To adjust the weights during training

Answer: c) To introduce non-linearity to the model

Question 71: What is the primary advantage of using Convolutional Neural Networks (CNNs) for image recognition tasks?

- a) Faster training time compared to traditional neural networks
- b) Ability to handle text data effectively
- c) Ability to capture spatial relationships in the image
- d) Higher accuracy on numerical datasets

Answer: c) Ability to capture spatial relationships in the image

Question 72: What is the main difference between a fully connected layer and a convolutional layer in a CNN?

- a) The fully connected layer has more neurons than the convolutional layer.
- b) The fully connected layer is used for image processing, while the convolutional layer is used for text processing.
- c) The fully connected layer applies convolutional operations to the input.
- d) The fully connected layer connects every neuron to every neuron in the previous layer.

Answer: d) The fully connected layer connects every neuron to every neuron in the previous layer.

Question 73: What is the primary purpose of using Recurrent Neural Networks (RNNs)?

- a) To handle sequential or time-series data
- b) To classify images
- c) To perform clustering
- d) To reduce the dimensionality of data

Answer: a) To handle sequential or time-series data

Question 74: What is the key characteristic of a Long Short-Term Memory (LSTM) cell in an RNN?

- a) It can process sequences in parallel.
- b) It can only handle short sequences of data.
- c) It uses a single layer for sequential processing.
- d) It can retain information over long time intervals.

Answer: d) It can retain information over long time intervals.

Question 75: What is the main goal of Natural Language Processing (NLP)?

- a) To analyze and visualize data
- b) To build and train neural networks
- c) To understand and generate human language
- d) To perform feature engineering

Answer: c) To understand and generate human language

Question 76: Which NLP task involves determining the sentiment (positive, negative, neutral) of a given text?

- a) Named Entity Recognition (NER)
- b) Part-of-Speech Tagging (POS)
- c) Sentiment Analysis
- d) Text Summarization

Answer: c) Sentiment Analysis

Question 77: Which IBM Watson service provides language translation capabilities?

- a) IBM Watson Assistant
- b) IBM Watson Discovery
- c) IBM Watson Language Translator
- d) IBM Watson Studio

Answer: c) IBM Watson Language Translator

Question 78: What is the primary purpose of using IBM Watson Studio in data science projects?

- a) To perform natural language processing tasks
- b) To build and train machine learning models
- c) To create conversational chatbots
- d) To store and manage data in the cloud

Answer: b) To build and train machine learning models

Question 79: What is the main benefit of deploying machine learning models on the cloud?

- a) Reduced model accuracy
- b) Lower cost of model deployment
- c) Limited access to the model for end-users
- d) Scalability and ease of access for end-users

Answer: d) Scalability and ease of access for end-users

Question 80: Which deployment method is suitable for machine learning models that require real-time predictions and low latency?

- a) Cloud Deployment
- b) On-Premise Deployment
- c) Edge Deployment
- d) Web Deployment

Answer: c) Edge Deployment

Question 81: What is the purpose of exploratory data analysis in data science?

- a) To preprocess and clean the data
- b) To build predictive models
- c) To analyze the historical data
- d) To gain insights and discover patterns in the data

Answer: d) To gain insights and discover patterns in the data

Question 82: Which step in the data science process involves identifying the data sources and collecting the data?

- a) Data Exploration
- b) Data Analysis
- c) Data Cleaning
- d) Data Collection

Answer: d) Data Collection

Question 83: Which of the following is NOT a valid Python data type?

- a) Integer
- b) String
- c) Dictionary
- d) Function

Answer: d) Function

Module 2: Python Basics

Question 84: What is the output of the following Python code snippet?

```
python
Copy code
x = [1, 2, 3]
y = x
```

x[0] = 4

print(y)

a) [1, 2, 3]

b) [4, 2, 3]

c) [1, 2, 3, 4]

d) [4, 2, 3, 4]

Answer: d) [4, 2, 3, 4]

Question 85: What is the purpose of data imputation in data wrangling?

a) To remove outliers from the dataset

b) To handle missing data

c) To scale numerical features to a similar range

d) To encode categorical variables

Answer: b) To handle missing data

Question 86: Which data wrangling technique is used to transform categorical variables into numerical representations?

a) Data Imputation

b) Data Normalization

c) Data Encoding

d) Data Scaling

Answer: c) Data Encoding

Question 87: In regression, the dependent variable is:

a) Categorical

b) Continuous

c) Discrete

d) Binary

Answer: b) Continuous

Question 88: What is the purpose of the coefficient of determination (R-squared) in regression?

a) To measure the variance of the independent variable

b) To measure the correlation between independent and dependent variables

c) To assess the goodness of fit of the regression model

d) To calculate the standard error of the regression coefficients

Answer: c) To assess the goodness of fit of the regression model

Question 89: Which algorithm is commonly used for binary classification tasks?

a) k-Nearest Neighbors (k-NN)

b) Random Forest

c) Support Vector Machine (SVM)

d) K-Means Clustering

Answer: c) Support Vector Machine (SVM)

Question 90: What is the output of the logistic function (sigmoid) used in logistic regression?

a) Positive integers

b) Negative integers

c) Real numbers between 0 and 1

d) Binary values (0 or 1)

Answer: c) Real numbers between 0 and 1

Question 91: What is the role of hyperparameters in machine learning algorithms?

a) To control the model complexity

- b) To represent the target variable
  - c) To identify outliers in the data
  - d) To perform feature selection
- Answer: a) To control the model complexity

Question 92: Which technique aims to find the optimal hyperparameters by trying different combinations exhaustively?

- a) Grid Search
  - b) Random Search
  - c) Gradient Descent
  - d) K-Means Clustering
- Answer: a) Grid Search

Question 93: What is the primary goal of unsupervised learning?

- a) To make predictions based on historical data
  - b) To classify data into different categories
  - c) To find patterns and structures in the data
  - d) To measure the performance of a model
- Answer: c) To find patterns and structures in the data

Question 94: Which unsupervised learning algorithm is commonly used for dimensionality reduction?

- a) Decision Tree
  - b) Principal Component Analysis (PCA)
  - c) Naive Bayes
  - d) Linear Regression
- Answer: b) Principal Component Analysis (PCA)

Question 95: In a binary classification problem, which metric is used to evaluate the model's performance when correctly identifying positive cases?

- a) Precision
  - b) Recall
  - c) F1-score
  - d) Accuracy
- Answer: b) Recall

Question 96: What is the formula for calculating the F1-score using precision and recall?

- a)  $F1\text{-score} = (\text{Precision} + \text{Recall}) / 2$
  - b)  $F1\text{-score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$
  - c)  $F1\text{-score} = \text{Precision} - \text{Recall}$
  - d)  $F1\text{-score} = \text{Precision} * \text{Recall}$
- Answer: b)  $F1\text{-score} = 2 * (\text{Precision} * \text{Recall}) / (\text{Precision} + \text{Recall})$

Question 97: What is the purpose of IBM Watson Assistant?

- a) To analyze and visualize data
  - b) To build and deploy machine learning models
  - c) To provide a conversational interface for applications
  - d) To perform natural language processing tasks
- Answer: c) To provide a conversational interface for applications

Question 98: Which IBM Watson service is used to analyze unstructured text data and extract insights from it?

- a) IBM Watson Assistant
- b) IBM Watson Discovery
- c) IBM Watson Studio

d) IBM Watson Machine Learning

Answer: b) IBM Watson Discovery

Question 99: What is the role of a REST API in machine learning model deployment?

a) To convert the model into a human-readable format

b) To handle data preprocessing before model deployment

c) To enable communication between the model and other software applications

d) To generate visualizations of the model's predictions

Answer: c) To enable communication between the model and other software applications

Question 100: Which deployment method is suitable for machine learning models that require real-time predictions and low latency?

a) Cloud Deployment

b) On-Premise Deployment

c) Edge Deployment

d) Web Deployment

Answer: c) Edge Deployment